

The HUMAN Digest 1.0

May 2024



Greetings. In this issue you will find:

1. Welcome from the HUMAN Digest
2. AI and the humanities
3. General news
4. Useful terms
5. People to know
6. Visualizing AI
7. Just for fun

Welcome to the HUMAN Digest

Welcome to the HUMAN Digest, a newsletter for Lake Forest College faculty and staff from the HUMAN grant. This newsletter curates select AI-related information that we believe will interest our community.

We recognize that there is a constant flow of information about AI, and it can be challenging to know what may be worth your time and attention. There are *many* articles, videos, newsletters, email solicitations, software programs, and social media posts. The HUMAN Digest will have classroom content but does not aspire to myopically focus on questions of academic honesty and AI. Rather, we are interested in innovative uses of this technology.

We aim to provide a curated selection that speaks to the work of the HUMAN grant and also interests our community. We anticipate publishing the HUMAN digest a few times between now and August 2024 and then settling into a regular pattern—likely every two weeks.

The content is assembled by Justin Kee, editor of the HUMAN Digest and technical advisor to the grant. Justin will join us as a full-time faculty member for 24-25, teaching a FIYS course called Digital Dawn: Humanity, Cyberspace and AI , two sections of The Humanistic Ethics of AI and working on the technical aspects of the project. He collaborates with HUMAN PI, Davis Schneiderman, and the HUMAN Steering Committee: Daw-Nay Evans, Kim Hazlett, Brian McCammack, David Sanchez-Burr and Tessa Sermet.

We are sending this to FacStaff because of the broad interest in AI on our campus. At the same time we don't want to add to your unwanted emails. If you prefer to not receive the HUMAN Digest in the future, please feel free to delete, skip or junk.

The HUMAN Digest welcomes feedback at kee@lakeforest.edu

Justin Kee, editor-in-chief

Davis Schneiderman, Principal Investigator, [HUMAN](#) grant

AI and the Humanities

Arizona State University hosts “hackathon” combining humanities and computer science students

At ASU, professors are [teaching](#) humanities students to lead with prompt-engineering skills combined with critical thinking.

<https://tech.asu.edu/features/ai-hackathon>

Kim Albrecht seeks to map ChatGPT’s inner “black box”

At Harvards metaLAB, Professor Albrecht explores how various domains of human knowledge are [represented](#) inside ChatGPT and what biases are present, whether in the training data or with humans filtering and curating the data itself.

Ethan Mollick’s new book written in collaboration with AI

Ethan Mollick, associate professor of management at Wharton, has an upcoming book about living and working with AI, called [Co-Intelligence](#). Notably, he used language models to augment, rather than substitute, for his own writing style and the content of his work.

ChatGPT and plagiarism in the classroom

Using [writing prompts and reflective exercises](#) to encourage students to maintain academic honesty, despite the availability of AI tools like ChatGPT.

<https://christinakatopodis.com/pedagogy/whos-afraid-of-chatgpt-humanities-skills-beat-ai-in-the-classroom>

General news

The Evil Robot overlords have arrived, 1.0



Figure, is a company dedicated to developing advanced humanoid robots powered by AI to address social needs, such as labor shortages or filling unsafe jobs. Most of our built environment is human-scale and Figure believes the robots would be able to navigate and complete tasks similarly to humans.

Figure's initial robot, the Figure 01, uses OpenAI's language modeling technology to understand and implement spoken commands.

See <https://figure.ai>

Devin, a new AI software agent

[Devin](#), from Cognition Labs, is an automated agent that works alongside humans to build software and debug code, among other tasks.

Pika, a novel text-to-video platform

Founded by two Stanford PhD students, [Pika](#) offers text-to-video, image-to-video and video-to-video generation.

Useful terms

Large language model. A large language model (“LLM”) is an artificial intelligence model that processes and generates language. LLM operate by simulating several neural networks to encode and process language. Models are trained on enormous amounts of text from various sources, such as specific datasets or from the internet.

Tokens. Words from datasets are often “tokenized” before training a language model or running inference on that model to create a unique vocabulary for the model. Tokens can be entire words or sub-pieces of words or special characters. By working with smaller units of text, the model can process and generate language more accurately.

Parameters. In the context of a LLM, the size or complexity of the model is generally measured by the number of parameters, which represent a complex set of values for the various layers of the neural network. The base model of ChatGPT has approximately 175 billion parameters. As an illustration of the rapid growth of the GPT class of models, GPT-1 (2018) had 117 million parameters and GPT-2 (2019) has 1.5 billion parameters.

Inference. This refers to the phase where a trained model is used to make predictions or decisions based on new or unseen data. After a model has been trained on a dataset, learning some of the underlying patterns and relationships within that data, inference utilizes this learned information to predict outcomes for novel inputs. This is the process that occurs when a user “prompts” a model.

People to know

[Geoffrey Hinton](#), a leading computer scientist at the University of Toronto. He pioneered work on multilayered neural networks, as well as using neural networks for computer vision. Last May, Hinton resigned from his role at Google in order to be able to freely speak about his concerns with the development of AI. Go deeper, “[Why the Godfather of AI fears what he’s built](#)”

A visualization

ChatGPT is given a prompt, “[Intelligence is](#)” hundreds of times and the various completions from the model are recorded and graphed. By mousing over the sequences of words, one can

visualize how the previous word or sequence of words helps the model generate the next word. Note: site may take a few moments to load.

If you wish to go deeper, a lengthy [explanation](#) on how ChatGPT works from Stephen Wolfram.

Just for fun

Make an (annoying) [song](#). (With thanks to Rachel Whidden for the suggestion).

~

HUMANS IN THE LOOP.

Have a tip or suggestion for an upcoming issue of the HUMAN Digest? Write Justin Kee an email at kee@lakeforest.edu

~

